

Derivations

Computation of the posterior $p(r_t \mid x_{1:t+h})$

Starting with $h = 1$, the joint distribution $p(r_t, x_{1:t+1})$ is:

$$p(r_t, x_{1:t+1}) = \sum_{r_{t+1}} p(r_t, r_{t+1}, x_{1:t+1}) \quad (1)$$

$$= \sum_{r_{t+1}} p(r_t, x_{1:t}) p(r_{t+1}, x_{t+1} \mid r_t, x_{1:t}) \quad (2)$$

$$= p(r_t, x_{1:t}) \sum_{r_{t+1}} p(x_{t+1} \mid r_t, r_{t+1}, x_{1:t}) p(r_{t+1} \mid r_t, x_{1:t}) \quad (3)$$

$$= p(r_t, x_{1:t}) \sum_{r_{t+1}} p(x_{t+1} \mid r_t, x_{t+1-r_{t+1}:t}) p(r_{t+1} \mid r_t) \quad (4)$$

Explanations:

- $p(x_{t+1} \mid r_t, r_{t+1}, x_{1:t}) = p(x_{t+1} \mid r_t, x_{t+1-r_{t+1}:t})$, because knowing r_{t+1} just selects the previous observations w.r.t. which we condition on. If $r_{t+1} = 0$, this becomes $p(x_{t+1} \mid r_t, x_{t+1-r_{t+1}:t}) = p(x_{t+1})$, which is the prior for observing x_{t+1} .
- The transition probability $p(r_{t+1} \mid r_t, x_{1:t})$ from the run length at time t to the run length at time $t+1$ does not depend on the history of observations $x_{1:t}$. This is a model assumption in BOCD.

For $h = 2$, the joint distribution $p(r_t, x_{1:t+2})$ is:

$$p(r_t, x_{1:t+2}) = \sum_{r_{t+1}, r_{t+2}} p(r_t, r_{t+1}, r_{t+2}, x_{1:t+2}) \quad (5)$$

$$= \sum_{r_{t+1}, r_{t+2}} p(r_t, r_{t+1}, x_{1:t+1}) p(r_{t+2}, x_{t+2} \mid r_t, r_{t+1}, x_{1:t+1}) \quad (6)$$

$$= \sum_{r_{t+1}, r_{t+2}} p(r_t, r_{t+1}, x_{1:t+1}) p(x_{t+2} \mid r_t, r_{t+1}, r_{t+2}, x_{1:t+1}) p(r_{t+2} \mid r_t, r_{t+1}, x_{1:t+1}) \quad (7)$$

$$= \sum_{r_{t+1}, r_{t+2}} p(r_t, r_{t+1}, x_{1:t+1}) p(x_{t+2} \mid r_{t+1}, x_{t+2-r_{t+2}:t+1}) p(r_{t+2} \mid r_{t+1}) \quad (8)$$

$$= p(r_t, x_{1:t}) \sum_{r_{t+1}, r_{t+2}} p(x_{t+1} \mid r_t, x_{t+1-r_{t+1}:t}) p(r_{t+1} \mid r_t) p(x_{t+2} \mid r_{t+1}, x_{t+2-r_{t+2}:t+1}) p(r_{t+2} \mid r_{t+1}) \quad (9)$$

Explanations:

- The same reasoning as before applies to both, the predictive distribution as well as the transition probability.
- In the last step, the previous result was used.

By following this pattern, for arbitrary $h \geq 0$ the joint distribution $p(r_t, x_{1:t+h})$ becomes:

$$p(r_t, x_{1:t+h}) = p(r_t, x_{1:t}) \underbrace{\sum_{r_{t+1}, \dots, r_{t+h}} \prod_{m=1}^h [p(x_{t+m} \mid r_{t+m-1}, x_{t+m-r_{t+m}:t+m-1}) p(r_{t+m} \mid r_{t+m-1})]}_{=:M} \quad (10)$$

The factor M is understood to be one, if $h = 0$.

The joint distribution $p(r_t, x_{1:t+h})$ can therefore be computed by the product of the joint distribution $p(r_t, x_{1:t})$, which BOCD computes anyway, and a factor M , which depends on the predictive and

transition probabilities at future time steps. These objects are components of the BOCD algorithm, computed at every time step and reusable internally within the implementation.

The run length posterior can be computed from the joint distribution:

$$p(r_t|x_{1:t+h}) = \frac{p(r_t, x_{1:t+h})}{p(x_{1:t+h})} = \frac{p(r_t, x_{1:t+h})}{\sum_{r_t} p(r_t, x_{1:t+h})} \quad (11)$$

The factor M can be best understood as being the product of h matrices M_{t+m} . Let's, e.g., assume that $t = 1$. The first matrix, for which $m = 1$, can be written as:

$$M_2 = \underbrace{\begin{bmatrix} p(x_2) & p(x_2|x_1) & 0 \\ p(x_2) & 0 & p(x_2|x_1) \end{bmatrix}}_{p(x_2|r_1, x_{2-r_2:1})} \circ \underbrace{\begin{bmatrix} p(r_2=0|r_1=0) & p(r_2=1|r_1=0) & 0 \\ p(r_2=0|r_1=1) & 0 & p(r_2=2|r_1=1) \end{bmatrix}}_{p(r_2|r_1)}, \quad (12)$$

where \circ denotes element-wise multiplication. Similarly, for $m = 2$ we have:

$$M_3 = \underbrace{\begin{bmatrix} p(x_3) & p(x_3|x_1) & 0 & 0 \\ p(x_3) & 0 & p(x_3|x_{1:2}) & 0 \\ p(x_3) & 0 & 0 & p(x_3|x_{1:2}) \end{bmatrix}}_{p(x_3|r_2, x_{3-r_3:2})} \quad (13)$$

$$\circ \underbrace{\begin{bmatrix} p(r_3=0|r_2=0) & p(r_3=1|r_2=0) & 0 & 0 \\ p(r_3=0|r_2=1) & 0 & p(r_3=2|r_2=1) & 0 \\ p(r_3=0|r_2=2) & 0 & 0 & p(r_3=3|r_2=2) \end{bmatrix}}_{p(r_3|r_2)} \quad (14)$$

For all matrices M_{t+m} , r_{t+m-1} increments along the rows and r_{t+m} increments along the columns, both starting at $r_{t+m-1} = r_{t+m} = 0$ in the upper-left element. Elements for which neither the condition $r_{t+m} = 0$ nor $r_{t+m-1} + 1 = r_{t+m}$ is satisfied are zero.

If h was just two, then $M = M_2 M_3$ in this example. The joint distribution can be written as a vector:

$$\begin{bmatrix} p(r_1=0, x_{1:3}) \\ p(r_1=1, x_{1:3}) \end{bmatrix} = \begin{bmatrix} p(r_1=0, x_1) \\ p(r_1=1, x_1) \end{bmatrix} \circ M \quad (15)$$

It is possible to use the particular structure of each matrix M_{t+m} (which has only elements in the first column and first upper diagonal) to implement this matrix multiplication efficiently (see `bocd._log_matmul_fast()`).